Paired mean changes with exploratory 95% case-bootstrap intervals

Held-out TEST | 12 released cases | one training seed | roles not pooled

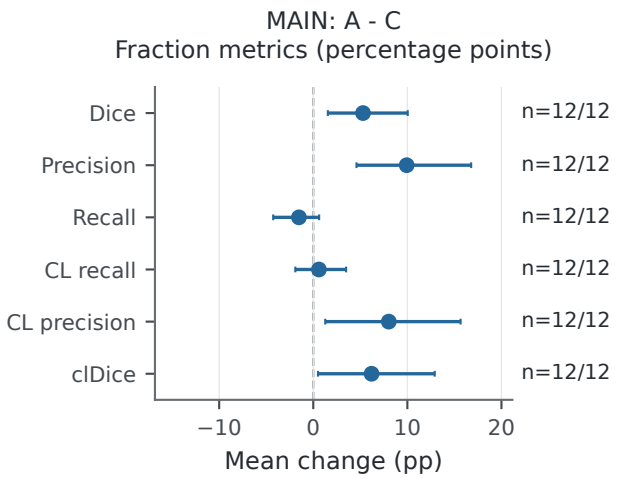

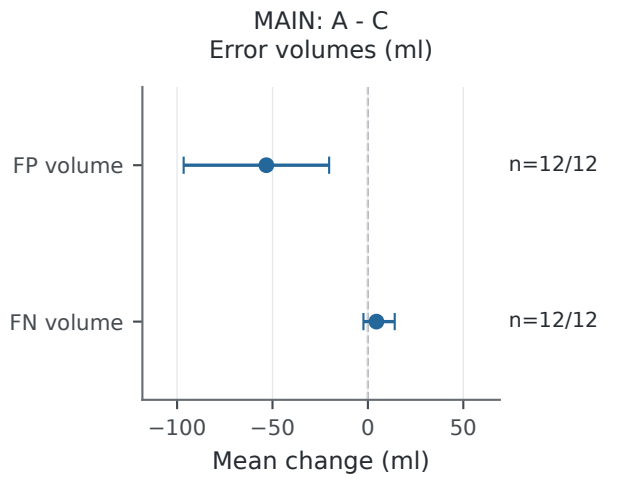

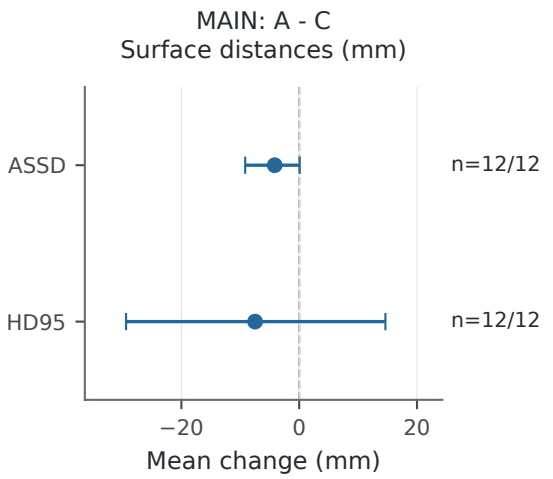

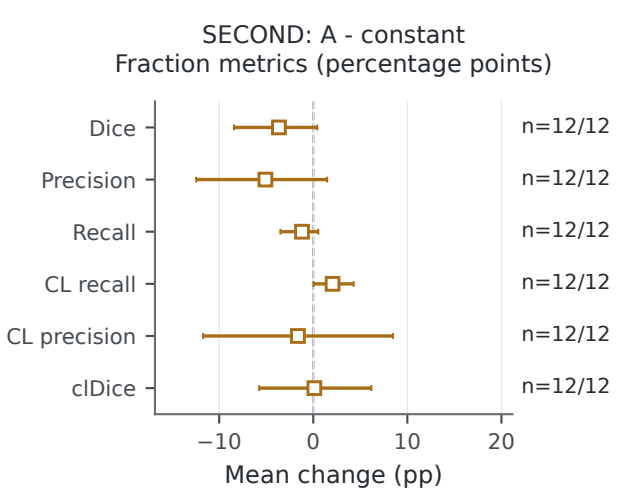

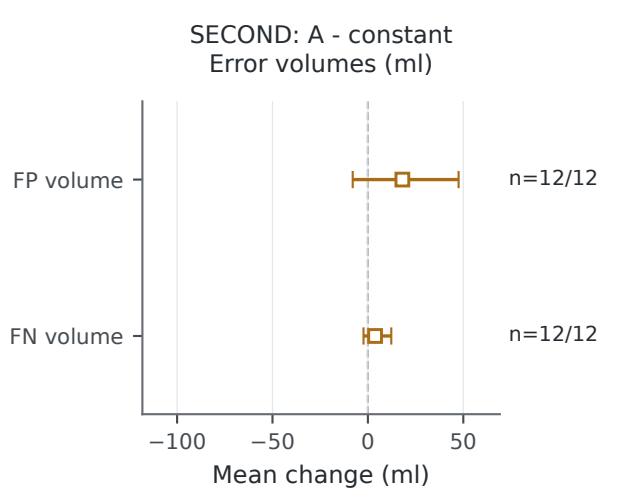

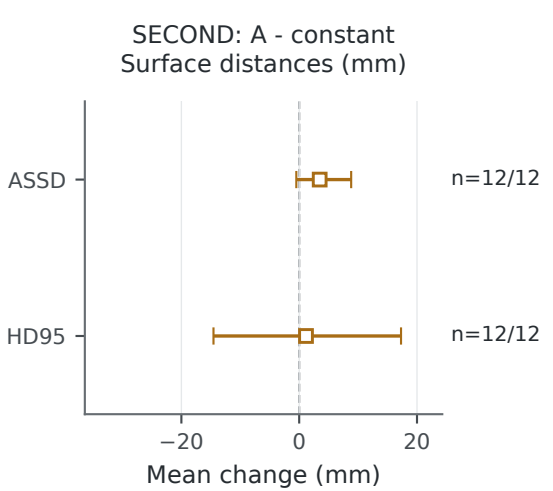

A minus its within-role comparator; effects are never pooled across roles. Larger fraction metrics are better; larger ml/mm errors are worse.
Intervals are read from the reviewed artifact, not recomputed. NA is undefined, not zero; n<2 intervals are degenerate. Case independence is unverified.

Published declaration/file hashes only; no author server or selection-timing attestation